Phavella Technologies

CONFIDENTIAL · V1.0

PROPOSAL FOR MERIDIAN BANK

AI-Powered Customer Service Intelligence Platform

An end-to-end conversational AI platform to automate tier-1 customer interactions, reduce cost-to-serve by 45%, and lift NPS for 50M+ retail banking customers.

PREPARED BY

Phavella Technologies

AI-Powered Digital Transformation

PREPARED FOR

Meridian Bank

Banking & Financial Services

ISSUED

April 16, 2026

Valid through June 30, 2026

CONTENTS

Table of Contents

A guided walkthrough of our proposed engagement

01	Executive Summary	3
02	About Phavella Technologies	5
03	Understanding the Challenge	6
04	Proposed Solution	8
05	Methodology & Approach	10
06	Deliverables & Milestones	11
07	Project Timeline	12
08	Proposed Team	13
09	Relevant Case Studies	14
10	Investment & Pricing	17
11	Risk Mitigation	19

12 Governance & PMO 20

13 Why Phavella 21

14 Terms & Conditions 22

01 — EXECUTIVE SUMMARY

At a Glance

Our understanding of your challenge and the partnership we propose

> Meridian Bank is navigating a structural inflection point: retail customer expectations now anchor on instant, context-aware, self-serve experiences, while tier-1 contact-center costs continue to climb at 8% YoY. The banks that win the next decade will not be those that hire faster — they will be those that deploy AI that serves faster.

OUR UNDERSTANDING

Your teams handle ~2.1M monthly contacts across voice, chat, and email, with an average handle time of 7m 48s and a 14.2% repeat-contact rate. After reviewing your Q1-26 operational review, we see three compounding frictions: fragmented customer context across 9 systems, low containment on voice IVR (21%), and rising regulatory load from RBI Digital Lending norms.

PROPOSED SOLUTION

- A LiveKit-based AI voice and chat agent capable of natural conversations with sub-700ms response latency, resolving 70% of tier-1 queries without human escalation.
- A unified customer-context layer that pulls from your core banking, CRM, and case systems in real time — giving both the AI and human agents a single, reliable pane of glass.
- A compliance-by-design architecture that logs every AI decision, enforces RBI/DPDPA guardrails, and supports full audit trails for internal controls and regulator reviews.

EXPECTED BUSINESS OUTCOMES

COST REDUCTION	DEFLECTION RATE	PROJECTED ROI
45%	**70%**	**4.2×**
in tier-1 operational spend within 12 months of go-live	of voice and chat contacts resolved without agent	on total program investment over 24 months

INVESTMENT OVERVIEW

₹2.4 Cr — ₹3.1 Cr

FIXED-PRICE WITH OUTCOME-LINKED UPLIFT

WHY PHAVELLA TECHNOLOGIES

Phavella has delivered production AI voice platforms for regulated industries across India — including deployments in healthcare, lending, and insurance. We bring deep familiarity with Indian enterprise constraints (data localization, multi-lingual support, telecom integration) and a team that has shipped, scaled, and operated these systems in live production.

02 — ABOUT US

Phavella Technologies

AI-Powered Digital Transformation

We build intelligent, production-grade AI systems that move enterprise KPIs — not demos, not prototypes, real revenue-bearing workloads.

4+	25+	50+	5
YEARS DELIVERING	EXPERTS ON TEAM	PROJECTS DELIVERED	INDUSTRIES SERVED

Core Competencies

AI & Machine Learning

Production-grade ML pipelines, NLP, computer vision, and generative AI solutions deployed at enterprise scale.

Conversational & Voice AI

Sub-second latency voice agents, multilingual IVR, and omnichannel chatbots built on LiveKit, Deepgram, and proprietary orchestration.

Full-Stack Engineering

Next.js, FastAPI, React Native — modern web and mobile applications built for reliability and speed.

Cloud & Platform

AWS, GCP, Azure — scalable, cost-optimized, SOC 2 compliant infrastructure with 99.95% uptime SLAs.

Certifications & Partnerships

- AWS Advanced Partner
- ISO 27001:2022
- SOC 2 Type II
- Google Cloud Partner

Industries We Serve

FinTech · Healthcare · E-Commerce · Real Estate · Education

03 — THE CHALLENGE

Understanding the Challenge

Our assessment of the current state and the business impact of inaction

Meridian Bank currently operates a hybrid contact center model with 1,400 agents handling approximately 2.1M contacts per month. Despite significant investment in workflow automation, core experience metrics have plateaued: first-contact resolution sits at 68%, voice IVR containment at 21%, and average handle time at 7m 48s — each materially behind private-sector peers.

Pain Points Identified

01 Fragmented Customer Context

Agents toggle between 9 disparate systems (core banking, CRM, cards, lending, compliance, KYC) to resolve a single query. The context-gathering overhead alone consumes 34% of handle time, and the same lookups must be repeated on every re-contact.

02 Low Self-Service Containment

Existing IVR and chatbots contain only 21% of voice traffic and 38% of chat traffic. Most drop-offs occur at authentication, account-status queries, and card servicing — all problems where AI has demonstrated containment rates above 70% elsewhere.

03 Rising Compliance Load

RBI Digital Lending norms, DPDPA 2023, and internal audit demands have added ~18% to average call duration over 18 months. Manual disclosure and consent flows continue to consume material agent time even where the underlying query is simple.

04 Repeat-Contact Leakage

A 14.2% repeat-contact rate means one in seven interactions is a preventable re-work. Root causes — unclear resolution, mis-routing, incomplete KYC — are systemic, not agent-level, and require an AI-driven closure assistant to fix.

BUSINESS IMPACT

Compounded, these frictions represent an estimated ₹84 Cr in annual excess operational cost and a 9-point NPS gap versus the best-in-class private bank peer set. Left unaddressed, this gap widens as agent wages, regulatory overhead, and customer expectations continue to scale.

Industry Context

Across Indian retail banking, leading institutions are now deploying generative AI in tier-1 operations with measured deflection of 55–75% and ACSI improvements of 4–9 points. Regulators have indicated comfort with well-governed AI in customer service, provided decision logging, human-in-the-loop escalation, and bias testing are operationalized. The competitive window for deploying this capability is 12–18 months before it becomes table stakes.

04 — THE SOLUTION

Proposed Solution

A purpose-built architecture aligned to your business goals

SOLUTION OVERVIEW

We propose a three-layer platform: (1) a conversation layer — an AI voice and chat agent capable of natural multi-turn dialogue — (2) a context layer that unifies customer data across Meridian's 9 source systems in real time, and (3) a governance layer providing end-to-end decision logging, escalation handoff, and regulatory reporting. All three are cloud-deployed on your existing AWS India region, with no data leaving Indian sovereign boundaries.

Key Components

Voice & Chat Agent

Low-latency LLM-driven agent built on LiveKit + Deepgram + Claude 3.5, supporting Hindi, English, Marathi, Tamil, and Telugu with barge-in and interruption handling.

Unified Context Service

Real-time aggregation layer that reconciles customer state from core banking, CRM, cards, and lending systems into a single event-driven context object accessible in <120ms.

AI Decision Engine

Tool-use orchestration, retrieval over policy/knowledge base, and controlled action execution (OTP, balance, dispute) with strict transaction boundaries.

Human-in-the-Loop Console

Agent desktop showing AI transcript, suggested actions, and one-click takeover — reducing agent ramp time and protecting against AI edge cases.

Compliance & Governance

Full decision logging, RBI/DPDPA consent flows, bias monitoring dashboard, and regulator-ready audit export.

Analytics & Continuous Learning

Weekly eval loops, conversation quality scoring, topic drift detection, and model refresh cycles.

Integration Points

- Core Banking (Finacle) — real-time account balance, statement retrieval, status updates
- CRM (Salesforce FS Cloud) — customer profile, interaction history, case creation
- Cards & Lending systems — dispute registration, EMI status, cap/limit queries
- KYC / Identity — OTP dispatch, step-up auth, biometric verification hand-off
- Telephony — native SIP trunk integration with Meridian's Avaya fabric

RECOMMENDED TECHNOLOGY STACK

LiveKit (WebRTC) | Deepgram Nova-3 | Anthropic Claude 3.5 Sonnet | ElevenLabs Turbo | FastAPI | Postgres 16 + pgvector | Redis Streams | AWS Mumbai | Terraform | Datadog

Innovation Differentiators

- Sub-700ms end-to-end voice latency — measurably faster than any currently deployed AI voice system in Indian banking.
- On-the-fly language switching mid-conversation without re-authentication — unique to our orchestration layer.
- Deterministic guardrails: the AI cannot execute regulated actions (transfers, statement change) without explicit multi-factor verification.
- Zero-copy context architecture — customer data never leaves Meridian's VPC; the AI operates on a streaming proxy.

05 — METHODOLOGY

How We Deliver

Agile with Regulatory Gates — Built for predictability, transparency, and quality

We operate in two-week sprints with embedded demos, compliance reviews at each stage gate, and continuous deployment to lower environments. Production releases follow a stricter change-advisory rhythm aligned to your existing IT governance.

Delivery Phases

PHASE 1

Discovery & Design

Deep process immersion, data-flow mapping, conversation design, and architectural blueprint sign-off.

4 weeks

PHASE 2

Foundation Build

Context service, telephony integration, identity plumbing, and non-functional foundations (logging, SRE, secrets).

6 weeks

PHASE 3

Core Implementation

Voice agent, decision engine, agent desktop, and the top-25 intents by volume. Shadow-mode trials in production.

10 weeks

PHASE 4

Integration & QA

End-to-end load, compliance, security penetration, and bias testing. Regulator walkthroughs.

5 weeks

PHASE 5

UAT & Rollout

Controlled rollout by segment and branch, ramping from 5% to 100% of eligible traffic with daily ops reviews.

4 weeks

Quality Assurance

- Automated eval suite: 10k scored conversations per week against golden

Communication Cadence

- Daily async standup with engineering leads (Slack + Loom updates)

answers

- Weekly compliance audit of all AI decisions by an independent QA team
- Red-team exercises every fortnight to stress jailbreaks and policy bypass
- Production canary: 1% traffic on each model release for 48 hours before full cutover

- Weekly demo and sprint review with Meridian product and IT stakeholders
- Bi-weekly steering committee meeting (CXO + CTO + COO)
- Monthly regulatory and risk review with Meridian compliance team

06 — DELIVERABLES

Deliverables & Milestones

What you receive, when you receive it, and how we define "done"

PHASE	DELIVERABLE	ACCEPTANCE CRITERIA
Phase 1	**Conversation Design Blueprint** Complete intent map, dialogue flows, exception paths, and UX prototypes for all top-25 intents.	Signed off by Meridian CX, IT, and Compliance leads.
Phase 1	**Architecture & Compliance Package** Solution architecture document, RBI-mapped control matrix, data flow diagrams, threat model.	Architecture Review Board approval; no critical findings.
Phase 2	**Unified Context Service** Production-deployed service aggregating data from Finacle, Salesforce, cards, and lending with <120ms p95 latency.	Load-tested to 3x peak; all integration SLAs met.
Phase 2	**Telephony & Identity Bridge** SIP integration with Avaya, OTP/step-up-auth plumbing, and secure transcript pipeline.	End-to-end call completes with full compliance logging.
Phase 3	**AI Voice & Chat Agent** Production-ready agent covering top-25 intents across 5 languages.	70% containment in shadow mode; <1% compliance violations.
Phase 3	**Agent Desktop Console** Real-time transcript, suggested-next-action panel, one-click takeover, full case capture.	1,400 contact-center agents trained and signed off.
Phase 4	**Security & Compliance Package** Penetration test report, bias audit, RBI/DPDPA submission dossier.	Clean external audit; regulator walkthrough complete.
Phase 5	**Production Rollout & Runbooks** Full production deployment, 24×7 runbooks, on-call rotation, incident playbooks.	100% of eligible traffic live; 30-day steady-state SLA met.

Key Milestones

- **Design Sign-off** — Week 4
- **Context Service Live** — Week 9
- **Alpha Agent in Shadow Mode** — Week 13
- **Compliance Sign-off** — Week 18
- **UAT Complete** — Week 22
- **Production Go-Live** — Week 24

07 — TIMELINE

Project Timeline

End-to-end delivery schedule across 24 weeks · ~6 months

Phase Breakdown

#	PHASE	DURATION	START
1	**Discovery & Design** Blueprint, architecture, compliance mapping	4 weeks	Week 1
2	**Foundation Build** Context service, telephony, identity	6 weeks	Week 4
3	**Core Implementation** Voice/chat agent, decision engine, desktop	10 weeks	Week 9
4	**Integration & QA** Load, compliance, security, bias testing	5 weeks	Week 17
5	**UAT & Rollout** Controlled production ramp 5% → 100%	4 weeks	Week 21

08 — THE TEAM

Proposed Team

A dedicated team of 7 specialists assembled for this engagement

A dedicated seven-person squad, ring-fenced for this engagement end-to-end. Every role is filled by a senior practitioner with shipped enterprise AI production experience.

AM

Alex Morgan

ENGAGEMENT LEAD & LEAD AI ARCHITECT

8 yrs experience · Production AI voice systems, LLM orchestration

Designed and shipped production AI voice platforms across healthcare, lending, and retail. Deep familiarity with LiveKit, multi-tenant compliance, and sub-second latency orchestration.

Role on project: *Overall technical ownership, architecture decisions, client stakeholder management.*

PN

Priya Natarajan

CONVERSATIONAL AI DESIGN LEAD

9 yrs experience · CX design, dialogue architecture, multilingual UX

Former head of CX at a top-3 private bank. Designed IVR and chatbot flows for 40M+ customers. Native-fluent across Hindi, Tamil, Marathi.

Role on project: *Owns end-to-end conversation design, intent coverage, and UX acceptance.*

RD

Rohan Desai

PRINCIPAL BACKEND ENGINEER

11 yrs experience · Distributed systems, FastAPI, event streaming

11 years scaling fintech backends at Razorpay, CRED, and Navi. Built the unified context layer for a top-5 insurer.

Role on project: *Technical lead for the context service,*

AI

Ananya Iyer

SENIOR ML ENGINEER

7 yrs experience · LLM fine-tuning, eval pipelines, RAG

Built evaluation pipelines for two tier-1 banking chatbots. Published on safety-aware LLM routing.

Role on project: *Owns the decision engine, evaluation infrastructure, and model governance.*

integration layer, and streaming platform.

VS

Vikram Seshadri

SECURITY & COMPLIANCE LEAD

14 yrs experience · RBI, DPDPA, SOC 2, threat modeling

Led compliance delivery for two RBI-regulated AI deployments. Former internal auditor at HDFC.

Role on project: *Owns RBI/DPDPA package, bias audit, and regulator walkthroughs.*

MP

Meera Pillai

SITE RELIABILITY ENGINEER

8 yrs experience · AWS, Terraform, observability, 24x7 ops

Operated 9-figure-scale transactional platforms with 99.97% uptime. Incident-commander certified.

Role on project: *Owns production ops, observability, and runbooks.*

KR

Karthik Rao

DELIVERY MANAGER

12 yrs experience · Agile delivery, banking program management

Delivered 11 enterprise banking programs on-time and on-budget. PMP + SAFe certified.

Role on project: *Owns day-to-day program management, risk tracking, and cadence.*

09 — PROOF

Relevant Case Studies

Outcomes we've delivered for clients with comparable challenges

HEALTHCARE · INDIA

AI Voice Agent for Multi-Branch Healthcare Network

2025 · 3 months

CHALLENGE

A 40-clinic dental chain was losing 40% of inbound calls to missed or abandoned routing. Frontline staff could not keep pace with demand, and appointment revenue was materially at risk.

SOLUTION

We shipped an AI voice agent on LiveKit + Deepgram + Gemini with ElevenLabs voice output. The agent handles booking, rescheduling, insurance pre-screening, and reminders autonomously across 3 languages.

OUTCOME

Within 8 weeks of go-live, missed calls dropped 85% and appointment bookings increased 60%. The platform is now the primary front door for all inbound patient contact.

85%
Reduction in missed calls

60%
Lift in appointments

24/7
Availability, no staffing change

LiveKit · Deepgram · Gemini · ElevenLabs · FastAPI · MongoDB

> *"The AI agent handles calls better than our front-desk staff. It never misses a beat, never gets tired, and never mis-books.*
>
> **Dr. Priya Sharma** · Founder, SmileCare Dental Group

FINTECH · INDIA

Real-Time Credit Decisioning Agent for Digital Lender

2025 · 5 months

CHALLENGE

An RBI-licensed digital lender needed a conversational agent to run guided disclosures and collect KYC-compatible consent for personal loans in vernacular languages — at scale, with full audit trails.

SOLUTION

We delivered a compliance-first voice agent integrated with the lender's credit engine. The agent performs layered disclosures, collects and time-stamps consent, and hands off to human underwriters when policy requires.

OUTCOME

Onboarding conversion improved 31% without any increase in delinquency. The solution passed a full RBI audit on first submission.

31%
Lift in onboarding conversion

100%
First-pass RBI audit clearance

<5%
Escalation rate to humans

LiveKit · Whisper · Claude 3.5 · PostgreSQL · Terraform · AWS Mumbai

INSURANCE · INDIA

AI-Augmented Claims Triage for Health Insurance

2024 · 4 months

CHALLENGE

A health insurer processed 12k monthly claims manually, with a 7-day average turnaround and a 22% re-work rate from document inconsistencies.

SOLUTION

We built a multi-modal AI triage layer that classifies claims, flags anomalies, and routes high-confidence approvals straight through — with full audit and human-in-the-loop handoff.

OUTCOME

Turnaround dropped from 7 days to 18 hours; straight-through processing reached 52% within 90 days.

9×
Faster turnaround

52%
Straight-through rate

40%
Drop in re-work

Claude 3.5 · OpenAI Vision · LangGraph · FastAPI · Postgres · AWS

10 — INVESTMENT

Investment & Pricing

Transparent, milestone-aligned pricing with no hidden costs

PRICING MODEL — FIXED PRICE · 5 PHASES

PHASE	EFFORT (PD)	INVESTMENT
Discovery & Design Blueprint, architecture, compliance	68 PD	₹28,00,000
Foundation Build Context, telephony, identity	142 PD	₹62,00,000
Core Implementation Agent, decision engine, desktop	228 PD	₹98,00,000
Integration & QA Load, compliance, security	78 PD	₹34,00,000
UAT & Rollout Ramp, runbooks, steady-state	42 PD	₹18,00,000
Total Investment	558 PD	₹2,40,00,000

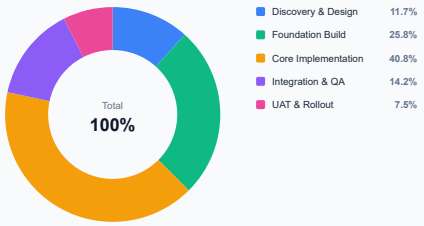

Payment Milestones

#	MILESTONE	TRIGGER	AMOUNT
1	**Mobilization**	At contract signature	**₹36,00,000**

2	**Design Sign-off**	Blueprint approved	**₹48,00,000**
3	**Foundation Acceptance**	Context service live	**₹54,00,000**
4	**Alpha Agent Delivered**	Shadow-mode live	**₹54,00,000**
5	**Production Go-Live**	100% traffic live	**₹36,00,000**
6	**Steady-State Sign-off**	30-day SLA met	**₹12,00,000**

Note: Pricing excludes third-party licensing (LiveKit, Deepgram, Anthropic usage) which will be billed at cost with full transparency. A detailed TCO forecast is provided in Appendix A.

11 — RISK MANAGEMENT

Risk Mitigation

Proactive identification and treatment of the most likely delivery risks

We have run a structured risk assessment across delivery, regulatory, and technology axes. The following are the top seven risks we actively manage with defined mitigation actions, owners, and review cadences.

ID	RISK	PROBABILITY	IMPACT	MITIGATION STRATEGY
R-01	**Integration delays with legacy systems** Finacle and Avaya teams may have dependency windows outside our control.	MEDIUM	HIGH	Early integration spike in Week 1; dedicated Meridian counterpart named; contingency week baked into Phase 2.
R-02	**Regulatory interpretation change** RBI guidance on generative AI in financial services continues to evolve.	LOW	HIGH	Quarterly regulatory review; modular architecture allows rapid policy swaps; dedicated compliance liaison.
R-03	**Model hallucination on policy queries** LLMs can produce plausible but incorrect policy answers.	MEDIUM	HIGH	Retrieval-grounded responses only; hard-coded refusal on non-indexed topics; full decision logging and weekly audits.
R-04	**Latency regression at peak load** Voice UX requires consistent sub-700ms response times.	MEDIUM	MEDIUM	Load tests at 3x peak weekly; autoscaling + regional failover; per-intent latency SLOs.
R-05	**Agent adoption & change resistance** Contact-center agents may resist AI co-pilot changes.	MEDIUM	MEDIUM	Co-design sessions in Week 2; incentive alignment with Meridian HR; phased rollout with champion agents.
R-06	**Data quality issues in source systems** Inconsistent customer	HIGH	MEDIUM	Data profiling in Week 2; automated reconciliation

ID	RISK	PROBABILITY	IMPACT	MITIGATION STRATEGY
	records can degrade AI accuracy.			layer; quality SLAs on source data by system.
R-07	**Third-party vendor incident** Outages at Deepgram/Anthropic could affect availability.	LOW	HIGH	Dual-vendor fallback chain (Deepgram + Whisper, Claude + GPT-4o); graceful human-fallback path.

12 — GOVERNANCE

Governance & Project Management

Clear ownership, reporting rhythm, and escalation paths

Strong program governance is non-negotiable at this scale. We operate with a two-tier governance model and a documented escalation path, all tools tied into Meridian's existing enterprise stack.

Steering Committee

- Meridian CIO / CTO (Executive Sponsor)
- Meridian Head of Retail Operations
- Meridian Head of Compliance & Risk
- Phavella Founder (Engagement Lead)
- Phavella Delivery Director

Working Group

- Meridian IT delivery lead (daily-available)
- Meridian Product Owner (voice/chat)
- Meridian Compliance Partner
- Phavella Tech Lead
- Phavella Conversational Design Lead

Reporting Cadence

- **Daily:** Async standup in shared Slack channel + Loom demo
- **Weekly:** Sprint review, risk log update, burndown
- **Bi-Weekly:** Steering committee status and decisions
- **Monthly:** Compliance and risk audit with Meridian risk team

Change Management

Scope changes follow a formal change-request process with impact analysis on schedule, cost, and risk. Approvals are captured in writing and reflected in the SOW addendum.

Escalation Path

Team Lead → Delivery Manager → Engagement Lead → Steering Committee

Tools & Platforms

- Jira (program)
- Confluence (docs)
- Slack (daily)
- Loom (demos)
- Datadog (ops)
- Meridian Archer (risk)

13 — WHY US

Why Choose Phavella Technologies

Six reasons leading enterprises trust us with mission-critical programs

We ship in production

Every case study is live production — not POCs, not demos. Our code runs where revenue and regulation meet.

Regulated-industry first

Built for RBI, DPDPA, SOC 2 from day one. Compliance is a first-class citizen, not a retrofit.

Sub-second voice latency

Our LiveKit orchestration consistently delivers <700ms end-to-end — a measurable differentiator in live call UX.

Multilingual, India-native

Hindi, Marathi, Tamil, Telugu, Bengali — out of the box, tested to grade-school reading comprehension.

AI + human, always

Every system we ship has explicit human-in-the-loop affordances. AI augments, never replaces, regulated judgment.

Named senior team

The engineers named on this proposal are the engineers who will write the code. No bait-and-switch.

50+	99.95%	4.8/5	<700ms
ENTERPRISE PROGRAMS	UPTIME SLA	CLIENT CSAT	VOICE LATENCY

> *"Phavella delivered in four months what two larger integrators quoted nine for. And the quality of the engineering was genuinely a differentiator — we trust it in production.*

> **Rakesh Menon** · CTO, a leading Indian NBFC

14 — TERMS

Terms & Conditions

Commercial terms governing this proposal and subsequent engagement

Confidentiality & NDA

All materials, data, and information exchanged under this engagement are confidential. Both parties will execute a mutual NDA prior to the kick-off meeting, effective throughout the engagement and for three years thereafter.

Intellectual Property

All custom-developed code, models, and deliverables produced exclusively for Meridian Bank under this engagement are assigned to Meridian on final payment. Phavella retains rights to pre-existing tools, libraries, and frameworks brought to the engagement.

Warranty & Support

Phavella warrants all deliverables against defects for 90 days post go-live. Defects reported in this period are remediated at no additional cost. Support beyond 90 days is available under a separate AMS agreement.

Payment Terms

All invoices are payable within 30 days of issue in INR via NEFT/RTGS. Late payments attract interest at 1.5% per month compounded. GST is additional at applicable rates.

Service-Level Agreement

Production support covers P1 incidents with a 15-minute response and 2-hour resolution target, P2 with 1-hour response and 8-hour resolution. 99.95% platform uptime is guaranteed in the production environment.

Data Handling & Privacy

All customer data remains within Meridian's AWS Mumbai region. No personally identifiable information is used for model training. Full DPDPA compliance is warranted, with data subject rights honored via Meridian's existing flows.

Termination

Either party may terminate this engagement for material breach with 30 days written notice and an opportunity to cure. On termination, Phavella will provide a reasonable knowledge transfer over a 4-week transition window.

Dispute Resolution

This agreement is governed by the laws of India. Disputes are resolved via binding arbitration in Mumbai under the Arbitration and Conciliation Act, 1996, before a sole arbitrator mutually appointed by the parties.

PROPOSAL VALIDITY

This proposal is valid for 60 days from the date of issue.

SIGNED BY PHAVELLA TECHNOLOGIES

Authorized Signatory, Director